

An Introduction to Monitoring Encrypted Network Traffic with Joy

Blake Anderson and David McGrew
Cisco Live US 2016 Workbench

Abstract: TLS encryption has become the standard form of Internet communication. In this session, we will demonstrate our open source project: Joy. It extends flow monitoring technologies by collecting much more detailed information about flows. This information can be used in conjunction with simple rules to detect obsolete cryptography on a network, or can be used by machine learning algorithms to detect malicious, encrypted flows. Both of these use cases will be highlighted.

Technical Level: Introductory
Technology: Open Source, Security
Solutions: Analytics, Threat Defense
Session Type: DevNet
Session Length: 45 min

Obtaining Joy

<https://github.com/davidmcgrew/joy>

make
install

```
sh$ sudo apt-get install build-essential python-dev python-numpy  
python-setuptools python-scipy libatlas-dev libatlas3-base  
sh$ sudo easy_install -U scikit-learn  
sh$ sudo apt-get install whois
```

Joy Components

pcap2flow
query.py
model.py
saltUI

Using pcap2flow to process pcaps into flow data files

```
sh$ pcap2flow bidir=1 http=1 dns=1 tls=1 dist=1 capture2.pcap >
capture2.gz
sh$ zless capture2.gz
{"version":"1.1","interface":"none","promisc":0,"daemon":0,"output":"
none","outputdir":"none","info":"none","count":0,"upload":"none","key
file":"none","retain":0,"bidir":1,"num_pkts":0,"type":1,"zeros":0,"di
st":1,"cdist":"none","entropy":0,"wht":0,"hd":0,"tls":1,"classify":0,
"idp":0,"dns":1,"exe":0,"anon":"none","useranon":"none","bpf":"none",
"verbosity":0}
{"sa":"10.0.2.15","da":"23.56.181.48","pr":6,"sp":43286,"dp":80,"ob":
12834,"op":93,"ib":173193,"ip":163,"ts":1465315516.880170,"te":146531
5547.566770,"ottl":64,"ittl":64,"otcp_win":14600,"itcp_win":65535,"ot
cp_syn":40,"otcp_nop":1,"otcp_mss":1460,"itcp_mss":1460,"otcp_wscale"
:7,"otcp_sack":1,"otcp_tstamp":1,"packets":[{"b":299,"dir": ">","ipt":
19},{ "b":1248,"dir": "<","ipt":86}, ...
```

The first line is metadata that describes the options used to convert the pcap into JSON. The second line is a JSON description of a flow:

```
{
  "sa":"10.0.2.15",           # source address
  "da":"23.56.181.48",       # destination address
  "pr":6,                     # protocol (TCP)
  "sp":43286,                 # source port
  "dp":80,                    # destination port (HTTP)
  "ob":12834,                 # number outbound bytes
  "op":93,                    # number outbound packets
  "ib":173193,                # number inbound bytes
  "ip":163,                   # number inbound packets
  "ts":1465315516.880170,     # time start (seconds since epoch)
  "te":1465315547.566770,     # time end (seconds since epoch)
  "ottl":64,                  # outbound IP TTL
  "ittl":64,                  # inbound IP TTL
  "otcp_win":14600,
  "itcp_win":65535,
  "otcp_syn":40,
  "otcp_nop":1,
  "otcp_mss":1460,
  "itcp_mss":1460,
  "otcp_wscale":7,
  "otcp_sack":1,
  "otcp_tstamp":1,
  "packets":[
    {"b":299,"dir": ">","ipt":19},
    {"b":1248,"dir": "<","ipt":86},
    ...
  ]
}
```

Using query to process JSON flow data files

We can get output that looks like Netflow:

```
sh$ query.py capture2.gz --summary | less
source address  destination address  prot  sport  dport  obytes  opkts  ibytes  ipkts  date       time       seconds
10.0.2.15      172.217.1.196       6     51932  443    710     8      311     9     2016-06-07 22:40:29 59.019
10.0.2.15      4.31.198.44         6     36350  80     1262    17     11068   16    2016-06-07 22:41:16 33.277
10.0.2.15      64.102.6.247        17    44385  53     32       1      214     1     2016-06-07 22:40:29 0.0
10.0.2.15      64.102.6.247        17    57384  53     52       2      900     2     2016-06-07 22:40:58 0.0
10.0.2.15      4.31.198.44         6     36344  80     1095    18     23760   27    2016-06-07 22:40:58 16.05
10.0.2.15      64.102.6.247        17    60086  53     26       1      444     1     2016-06-07 22:40:58 0.0
```

But we can also get a lot more information:

```
sh$ query.py capture2.gz | less

{
  "itcp_mss": 1460,
  "ip": 9,
  "i_probable_os": "FreeBSD / OS X",
  "ib": 311,
  "pr": 6,
  "otcp_syn": 40,
  "otcp_win": 14600,
  "ts": 1465339229.163859,
  "ottl": 64,
  "te": 1465339288.182719,
  "otcp_nop": 1,
  "itcp_win": 65535,
  "otcp_mss": 1460,
  "otcp_sack": 1,
  "da": "172.217.1.196",
  "otcp_wscale": 7,
  "dp": 443,
  "ittl": 64,
  "otcp_tstamp": 1,
  "sp": 51932,
  "packets": [
    {
      "b": 517,
      "ipt": 32,
      "dir": ">"
    },
    {
      "b": 168,
      "ipt": 28,
      "dir": "<"
    },
    {
      "b": 87,
      "ipt": 0,
      "dir": ">"
    },
    {
      "b": 65,
      "ipt": 0,

```

The `--select` option will select certain fields to be printed.

```
query.py capture2.gz --select "sa, da, dp"
```

The `--where` option will filter the flows to meet certain criteria.

```
query.py capture2.gz --select "sa, da, dp" --where "pr=17"
```

High entropy flows can be identified by using the Byte Distribution

```
sh$ query.py capture2.gz --select "sa, da, dp, bd" --where
"entropy(bd)>7.93"
{
  "select": [
    {
      "sa": 10.0.2.15 , "da": 4.31.198.44 , "dp": 80 , "bd":
[182, 40, 30, 39, 36, 31, 34, 44, 27, 33, 39, 39, 32, 39, 36, 49, 25,
34, 31, 28, 29, 40, 37, 37, 40, 40, 36, 38, 40, 38, 41, 57, 40, 27,
27, 44, 42, 33, 32, 40, 29, 34, 38, 51, 36, 46, 46, 49, 34, 42, 42,
38, 41, 49, 37, 36, 42, 46, 45, 36, 46, 47, 49, 60, 22, 30, 28, 36,
21, 34, 41, 46, 37, 47, 35, 43, 28, 41, 37, 58, 27, 34, 37, 42, 32,
40, 39, 52, 39, 40, 51, 50, 38, 42, 45, 62, 28, 38, 34, 48, 37, 55,
31, 46, 36, 60, 43, 59, 34, 50, 50, 57, 35, 43, 40, 47, 44, 50, 38,
47, 55, 43, 40, 40, 51, 44, 53, 61, 24, 30, 38, 35, 26, 36, 31, 39,
31, 35, 35, 34, 44, 40, 43, 51, 30, 38, 43, 33, 28, 41, 44, 41, 23,
29, 45, 37, 37, 36, 43, 52, 22, 32, 30, 42, 38, 44, 34, 50, 36, 44,
35, 46, 37, 47, 42, 60, 28, 37, 36, 36, 50, 51, 43, 54, 36, 44, 45,
40, 43, 39, 54, 57, 32, 42, 35, 45, 45, 39, 45, 53, 33, 37, 34, 42,
32, 41, 37, 53, 29, 39, 45, 46, 40, 38, 44, 58, 32, 39, 50, 45, 38,
43, 42, 47, 46, 53, 57, 54, 35, 44, 35, 41, 31, 42, 40, 46, 36, 49,
41, 45, 67, 67, 39, 43, 41, 44, 54, 42, 84, 44, 44, 55, 70, 52, 63,
68] }
    ]
  }
}
```

Looking at TLS

```

sh$ query.py capture2.gz --where "dp=443" | less
"tls": {
  "tls_irandom":
"57574d7d1e6a95a0304ccc2d71bb14765c2bd8a3cd966039bae034329471e462",
  "tls_iv": 5,
  "tls_ov": 5,
  "SNI": [
    "www.google.com"
  ],
  "srlt": [
    {
      "b": 512,
      "tp": "22:1",
      "ipt": 0,
      "dir": ">"
    },
    "tls_isid":
"5ff9b9c4356c4c102f3139daa4f58cfa968c5ff49b201fab81b287853e1c8c3a",
    "tls_osid":
"5ff9b9c4356c4c102f3139daa4f58cfa968c5ff49b201fab81b287853e1c8c3a",
    "tls_ext": [
      {
        "data": "00",
        "length": 1,
        "type": "ff01"
      },
      "tls_orandom":
"4e6fbeb21d3549c945fd52f17dc4190f195561117b970866dfd8651225691e59",
      "cs": [
        "c02b",
        "c02f",
        "c00a",
        "c009",
        "c013",
        "c014",
        "c012",
        "c007",
        "c011",
        "0033",
        "0032",
        "0045",
        "0039",
        "0038",
        "0088",
        "0016",
        "002f",
        "0041",
        "0035",
        "0084",
        "000a",
        "0005",
        "0004"
      ],
      "s_tls_ext": [
        {
          "data": "00",
          "length": 1,
          "type": "ff01"
        },
        {
          "data": "02683208737064792f332e3108687474702f312e31",
          "length": 21,
          "type": "3374"
        },
        {
          "data": "0100",
          "length": 2,
          "type": "000b"
        }
      ],
      "scs": "c02f"
    }
  ]
}

```

Looking at TLS security levels

```
sh$ query.py capture2.gz --where "dp=443" --select "sa, da,
seclevel(tls)"
{
  "name": [
    { "sa": "10.0.2.15" , "da": "172.217.1.196" ,
    "seclevel(tls)": "recommended" } ,
    { "sa": "10.0.2.15" , "da": "172.217.1.196" }
  ]
}
```

Identify malware flows based on SPLT and BD

The built-in classifier can detect malware based on its Sequence of Packet Lengths and Times (SPLT) behavior and its Byte Distribution:

```
sh$ pcap2flow bidir=1 dist=1 classify=1 capture2.pcap > capture2.gz
sh$ query.py capture2.gz --select "da, dp, p_malware" --where
"p_malware > 0.01"
{ "select": [
  { "da": 4.31.198.44 , "dp": 80 , "p_malware": 0.038026 } ,
  { "da": 4.31.198.44 , "dp": 80 , "p_malware": 0.035637 } ,
  { "da": 4.31.198.44 , "dp": 80 , "p_malware": 0.011685 }
]
```

Example showing a similar run on a malicious PCAP:

```
sh$ pcap2flow bidir=1 dist=1 classify=1
133d773a8ca64c24bd81b594cf5240dd.pcap > malware.gz
sh$ query.py malware.gz --select "da, dp, p_malware" --where
"p_malware > 0.99"
{
  "name": [
    { "da": 86.59.21.38, "dp": 443 , "p_malware": 0.997 } ,
    { "da": 108.61.179.216, "dp": 443 , "p_malware": 0.995 } ,
    { "da": 5.9.123.81, "dp": 9001 , "p_malware": 0.999 } ,
    { "da": 109.238.6.25, "dp": 443 , "p_malware": 0.992 } ,
    { "da": 176.31.159.231, "dp": 443 , "p_malware": 0.998 } ,
    ...
  ]
}
```

Making a classifier

We can use `model.py` to learn a logistic regression classifiers from two data directories, one containing malicious output and one containing benign output. The arguments to `model.py` are:

```
-p POS_DIR, --pos_dir POS_DIR      Directory of Positive Examples (JSON Format)
-n NEG_DIR, --neg_dir NEG_DIR      Directory of Negative Examples (JSON Format)
-m, --meta                          Parse Metadata Information
-l, --lengths                       Parse Packet Size Information
-t, --times                         Parse Inter-packet Time Information
-d, --dist                          Parse Byte Distribution Information
-o OUTPUT, --output OUTPUT         Output file for parameters
```

Joy uses two classifiers, one with only metadata and packet lengths/times, and one that also contains the byte distribution. To generate the parameter file that does use the byte distribution, we run:

```
sh$ pcap2flow bidir=1 dist=1 malware/*.pcap >
malware_train/malware.gz
sh$ pcap2flow bidir=1 dist=1 benign/*.pcap > benign_train/benign.gz
sh$ python ../joy/saltUI/model.py -m -l -t -p malware_train/ -n
benign_train/ -o params.txt

Features Used:
    Metadata                (7)
    Packet Lengths          (100)
    Packet Times             (100)
Total Features: 207

non-zero parameters:      21
```

and then we generate the parameters that do use the byte distribution:

```
sh$ python ../joy/saltUI/model.py -m -l -t -d -p malware_train/ -n
benign_train/ -o params_bd.txt
Num Positive:    599
Num Negative:    293

Features Used:
    Metadata                (7)
    Packet Lengths          (100)
    Packet Times             (100)
    Byte Distribution        (256)
Total Features: 463

non-zero parameters:      24
```

We can now visualize these classifiers by copying the `params.txt` and `params_bd.txt` files to the `saltUI` directory. Edit the `laui.cfg` file by replacing the line:

```
classifier Malware      logreg  logreg parameters.txt  logreg parameters_bd.txt
```

with:

```
classifier Malware      logreg  params.txt  params_bd.txt
```

In the `saltUI` directory, start the UI with:

```
python server.py
```

And point your browser to <http://localhost:8080/home>, and click on Local Analytics -> Upload JSON File:

Select a JSON file:

[Browse ...](#)[Start upload](#)

Upload a malicious, processed pcap, and the results using the new classifiers will be displayed:

Number of flows classified: 589

P(Malware) ▾	P(TLS) ▾	P(CKL) ▾	Source Address ▾	Dest.Address ▾	Source Port ▾	Dest. Port ▾	Inbound Packets ▾	Outbound Packets ▾
1.0	1.0	-1.0	172.16.45.20	104.237.132.39	1034	443	15	11
1.0	1.0	-1.0	172.16.45.20	104.237.132.39	1035	443	15	11
1.0	1.0	-1.0	172.16.45.20	104.237.132.39	1036	443	15	11
1.0	1.0	-1.0	172.16.45.20	104.237.132.39	1037	443	15	11
1.0	1.0	-1.0	172.16.45.20	104.237.132.39	1038	443	15	11
1.0	1.0	-1.0	172.16.45.20	104.237.132.39	1039	443	15	11
1.0	1.0	-1.0	172.16.45.20	104.237.132.39	1040	443	15	11
1.0	1.0	-1.0	172.16.45.20	104.237.132.39	1041	443	15	11
1.0	1.0	-1.0	172.16.45.20	104.237.132.39	1042	443	15	11
1.0	1.0	-1.0	172.16.45.20	104.237.132.39	1044	443	15	11
1.0	1.0	-1.0	172.16.45.20	104.237.132.39	1045	443	15	11